

Working with standard optimization packages

Jonathan F.C. Herman


Colorado Center for Astrodynamic Research

University of Colorado at Boulder



Optimal Trajectories

- Personal Background
- Common elements
 - General description
 - Toy problem formulation
 - Scaling
 - Sparsity
 - Local optima
- Optimizer differences
 - SNOPT vs IPOPT vs MATLAB (fmincon, ...)
- Advanced example
- Project suggestion


- Spent my master's thesis writing a Sims-Flanagan type low-thrust trajectory optimization tool
 - Essentially a watered down version of JPL's MALTO
 - Interfaced with SNOPT for optimization (same as MALTO)
- Used for optimizing far-side, highly inclined solar observatory trajectories



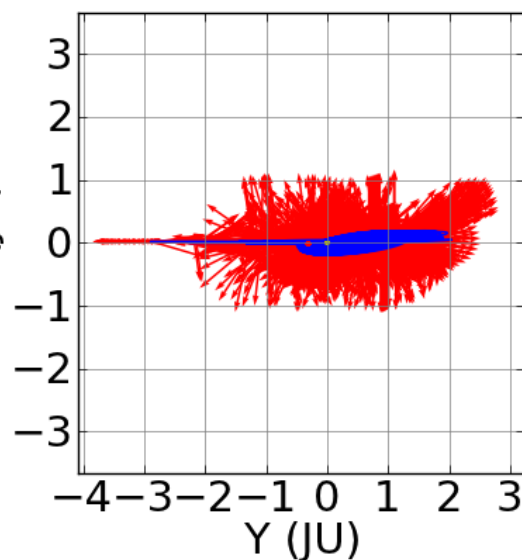

- 7.5 year time of flight
- Orbital period of ~ 1 year
- Final heliocentric inclination of 90 degrees
- $C3 = 0$, no gravity assists, all solar electric propulsion




- Realized that even though this code was very specific in its application, it applied very readily to all sorts of problems
 - Human missions to near Earth asteroids
 - Jovian moon tour (GTOC 6)
 - “Warp drive” studies
 - Human fly-by mission of Mars
 - Robotic asteroid tours (GTOC 7)
- Current activities
 - Retired my original low-thrust code
 - Helping out developing a much more powerful implementation of the same method (championed by Stijn de Smet)
 - Working on effective methods for much more complex problems



- GTOC6
- Scaling gone wrong...
- But it worked, and got us a very nice ranking!





- Spent about 3 years working with SNOPT for a variety of problems & problem formulations
 - And about a month with IPOPT
- Racked up a long list of do's and don'ts over that time
- Hoping to save you all some time by discussing them today
 - Very implementation focused
 - Plenty of published text on the theory behind these tools, easily found through google (or contact me)

Optimization tools

- At a high level, most optimization libraries are very similar
 - Enormous differences can exist in underlying theory / preferred interface / available options / actual performance
 - Besides those, there are some common features that, if not properly understood, can dominate the effects above
 - The interfaces are ultimately also very similar
- Before we get to some specific differences, its helpful to go over some of these potentially dominant features

- Parameters (1 x n)

$$x^L \leq x \leq x^U$$

- Cost function (1 x 1)

$$f(x)$$

- Constraints (m x 1)

$$g^L \leq g \leq g^U$$

- Cost derivatives (1 x n)

$$\nabla f(x) = \frac{\delta f}{dx}$$

- Jacobian (m x n)

$$\nabla g(x) = \frac{\delta g}{dx}$$

- For the purposes of this talk, it is okay to assume that the provided information (constraints, objective, derivatives) are used in a process similar to how a Newton solver is iteratively used to find a root
 - This is obviously an enormous simplification, but there are actually great similarities in these processes

$$X_1 = X_0 - \frac{f(X_0)}{f'(X_0)}$$

- Key take-away:
 - The solver will iteratively improve the optimization state, all the time computing the current value of the objective/constraints, and their derivatives, to determine the state for the next iteration

- Some solvers also allow the use of second order derivatives
 - Can be a useful feature, but can also be very memory intensive
 - It doesn't really change the aspects discussed today, so will largely be ignored here
- The formulation so far is very generic
 - As expected, smooth functions for the constraints/objective lead to better convergence in this approach
 - Adjoint equations need not be in user formulation
 - Since they are continuous functions, they could be though (and some trajectory optimization tools do this)
 - Typically though, these tools are used for direct optimization

- To give some concrete examples, we'll work with a toy problem
 - 3 parameters, 2 constraints
 - Designed for instructive/demonstrative value
 - It (probably?) doesn't solve anything useful
 - It should still have a name...

- To give some concrete examples, we'll work with a toy problem
 - 3 parameters, 2 constraints
 - Designed for instructive/demonstrative value
 - It (probably?) doesn't solve anything useful
 - It should still have a name...
- We will dub this "The POPSICLE-problem"

**The Pointless Optimization Problem
So I Can Learn Every-problem**

- Parameters (1 x 3)
- Cost function (1 x 1)
- Constraints (2 x 1)
- Cost derivatives (1 x 3)
- Jacobian (2 x 3)


$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}^T \leq \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}^T \leq \begin{bmatrix} 5 \\ 5 \\ 5 \end{bmatrix}^T$$

$$f(x) = x_0^2 + x_1^3 + x_2^2$$

$$\begin{bmatrix} 0 \\ 30 \end{bmatrix} \leq \begin{bmatrix} g_1 = x_0^3 + x_2 \\ g_2 = x_1^2 + x_2^3 \end{bmatrix} \leq \begin{bmatrix} 2.5 \\ 100 \end{bmatrix}$$

$$\nabla f(x) = [2x_0 \quad 3x_1^2 \quad 2x_2]$$

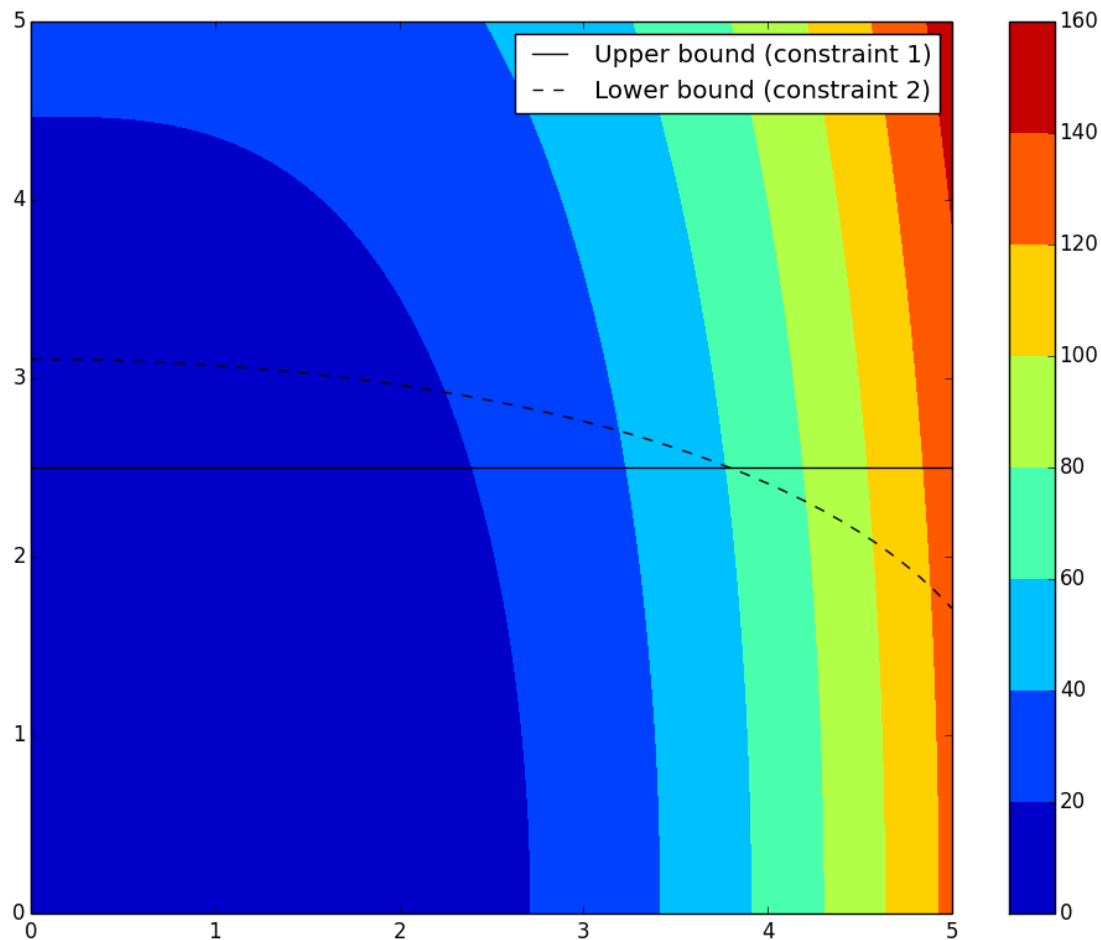
$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$




$$x_0 = 0$$

$$\begin{bmatrix} g_1 \\ g_2 \end{bmatrix} = \begin{bmatrix} x_0^3 + x_2 \\ x_1^2 + x_2^3 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ 30 \end{bmatrix} \leq \begin{bmatrix} g_1 \\ g_2 \end{bmatrix} \leq \begin{bmatrix} 2.5 \\ 100 \end{bmatrix}$$





The importance of scaling!

- Parameters (1 x 3)
- Cost function (1 x 1)
- Constraints (2 x 1)
- Cost derivatives (1 x 3)
- Jacobian (2 x 3)

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}^T \leq \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}^T \leq \begin{bmatrix} 5 \\ 5 \\ 5 \end{bmatrix}^T$$

$$f(x) = x_0^2 + x_1^3 + x_2^2$$

$$\begin{bmatrix} 0 \\ 30 \end{bmatrix} \leq \begin{bmatrix} g_1 = x_0^3 + x_2 \\ g_2 = x_1^2 + x_2^3 \end{bmatrix} \leq \begin{bmatrix} 2.5 \\ 100 \end{bmatrix}$$

$$\nabla f(x) = [2x_0 \quad 3x_1^2 \quad 2x_2]$$

$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$

- Parameters (1 x 3)

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}^T \leq \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}^T \leq \begin{bmatrix} 5 \\ 5 \cdot 10^7 \\ 5 \end{bmatrix}^T$$

- Cost function (1 x 1)

$$f(x) = x_0^2 + (x_1 \cdot 10^{-7})^3 + x_2^2$$

- Constraints (2 x 1)


$$\begin{bmatrix} 0 \\ 30 \end{bmatrix} \leq \begin{bmatrix} g_1 = x_0^3 + x_2 \\ g_2 = (x_1 \cdot 10^{-7})^2 + x_2^3 \end{bmatrix} \leq \begin{bmatrix} 2.5 \\ 100 \end{bmatrix}$$


- Cost derivatives (1 x 3)

$$\nabla f(x) = [2x_0 \quad 3x_1^2 \cdot (10^{-7})^3 \quad 2x_2]$$

- Jacobian (2 x 3)

$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 \cdot (10^{-7})^2 & 3x_2^2 \end{bmatrix}$$





- Clearly, this arbitrary scale factor of 10^{-7} is absurd, right?
 - Sure, except... $\frac{V_{Earth}}{AU} \approx \frac{30 \text{ km/s}}{150 \cdot 10^6 \text{ km}} = 2 \cdot 10^{-7}$
- **Scaling. Is. Everything.** It can (or rather; It will)...
 - Make a feasible solution appear infeasible
 - Make an infeasible solution appear feasible
 - Change runtimes by orders of magnitude
 - Determine whether or not the problem converges at all
 - Give you enormous headaches
- Scale appropriately!
 - **Everything** should be in the $[0,1]$ domain (or $[-0.5,0.5]$)
 - That means parameters, constraints, and objective function
 - You can deviate to weight certain aspects, **but tread carefully!**

- Most optimization problems have many derivatives that are always zero → Sparsity
- An important way to reduce memory usage
 - Which can also lead to improved runtimes
- Easy to show implementation with the POPSICLE-problem
 - Harder to show importance
- Will follow the sparse POPSICLE-problem with a real example

- Actual Jacobian

$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$

- Actual Jacobian

$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$

- Input of dense (full) Jacobian

$$\nabla g(x) = [3x_0^2 \quad 0 \quad 1 \quad 0 \quad 2x_1 \quad 3x_2^2]$$

$$iRow = [0 \quad 0 \quad 0 \quad 1 \quad 1 \quad 1]$$

$$jCol = [0 \quad 1 \quad 2 \quad 0 \quad 1 \quad 2]$$

- Actual Jacobian

$$\nabla g(x) = \begin{bmatrix} 3x_0^2 & 0 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$

- Input of dense (full) Jacobian

$$\nabla g(x) = [3x_0^2 \quad 0 \quad 1 \quad 0 \quad 2x_1 \quad 3x_2^2]$$

$$iRow = [0 \quad 0 \quad 0 \quad 1 \quad 1 \quad 1]$$

$$jCol = [0 \quad 1 \quad 2 \quad 0 \quad 1 \quad 2]$$

- Input of sparse Jacobian

$$\nabla g(x) = [3x_0^2 \quad 1 \quad 2x_1 \quad 3x_2^2]$$

$$iRow = [0 \quad 0 \quad 1 \quad 1]$$

$$jCol = [0 \quad 2 \quad 1 \quad 2]$$

- Example from my own research
- Low-thrust trajectory split into segments
 - Each segment contains a number of parameters & constraints
 - Sparsity grows (almost) linearly with problem size

Segments	100	1000	10 000
Non-zero Jacobian elements (%)	1.8	0.18	0.018
Full Jacobian (megabyte)	7.1	700	70 000
Sparse Jacobian (megabyte)	0.1	1.3	12.6

- Sparsity matters!

Local optima

- Parameters (1 x 3)

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}^T \leq \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}^T \leq \begin{bmatrix} 5 \\ 5 \\ 5 \end{bmatrix}^T$$

- Cost function (1 x 1)

$$f(x) = x_0 + 0.01x_1 + x_2(1.6 - 0.3x_2)$$

- Constraints (2 x 1)

$$\begin{bmatrix} -10 \\ 10 \end{bmatrix} \leq \begin{bmatrix} g_1 = -2x_1 + x_2 \\ g_2 = x_1^2 + x_2^3 \end{bmatrix} \leq \begin{bmatrix} 0 \\ 100 \end{bmatrix}$$


- Cost derivatives (1 x 3)

$$\nabla f(x) = [1 \quad 0.01 \quad 1.6 - 0.6x_2]$$

- Jacobian (2 x 3)


$$\nabla g(x) = \begin{bmatrix} 0 & -2 & 1 \\ 0 & 2x_1 & 3x_2^2 \end{bmatrix}$$


$$f(x) = x_0 + 0.01x_1 + x_2(1.6 - 0.3x_2)$$




$$\begin{bmatrix} g_1 = -2x_1 + x_2 \\ g_2 = x_1^2 + x_2^3 \end{bmatrix}$$

$$\begin{bmatrix} -10 \\ 10 \end{bmatrix} \leq \begin{bmatrix} g_1 \\ g_2 \end{bmatrix} \leq \begin{bmatrix} 0 \\ 100 \end{bmatrix}$$







(Switch to code)

- SNOPT
- IPOPT
- MATLAB (fmincon, ...)
- Many others exist
 - GPU-based
 - ...

- SNOPT (Sparse Nonlinear OPTimizer)
 - Written in Fortran 77 (Fortran 2003 version in development)
 - Interfaces exist for C, C++, Fortran, MATLAB, Python, ...
 - CU has an academic license to a fairly recent version
 - Sparsity notation a little more cumbersome than IPOPT
 - No parallel solving, and only first order derivatives
 - Performance generally excellent
 - Has become somewhat of an industry standard in optimization
 - Assumes/prefers you install on Linux...
 - Website: <http://ccom.ucsd.edu/~optimizers/>


- IPOPT (Interior Point OPTimizer)
 - Written in C++
 - Interfaces exist for C, C++, Fortran, MATLAB, Python, ...
 - Completely functional as open source distribution
 - Additional capability available through (academic) licenses
 - Very modular in design, essentially a wrapper for several solvers
 - Support for parallel solving & second order derivatives
 - Sparsity notation extremely straightforward
 - Seems to perform somewhat-to-significantly better than SNOPT
 - Very widely used (and seems to be growing)
 - EXCELLENT documentation & community support
 - Assumes/prefers you install on Linux...
 - Website: <http://www.coin-or.org/Ipopt/>

- fmincon (MATLAB)
 - Easy to get access to (it's in MATLAB...done)
 - MATLAB probably also has other solvers
 - No personal experience
 - Similarly in its setup/functionality to SNOPT/IPOPT
 - It does support sparse matrices
 - Might be perfect for a project, probably not for research
 - Website: <http://www.mathworks.com/help/optim/ug/fmincon.html>



- Instead of the user defining the Jacobian, some optimizers offer the functionality of finite differencing the objective & constraint functions
 - Generally worthless, since it comes at tremendous computational cost and is less accurate than analytical derivatives
 - It does allow for very quick implementations of problem formulations, so there is some value during prototyping
- Similarly, most optimizers can check the user-provided derivatives through finite differencing
 - This is **invaluable**, since it allows for (comparatively) easy identification of scaling errors, indexing errors, and 100 other errors you didn't know were even possible...

- It can be a chore to get SNOPT/IPOPT to compile
 - Especially if you are not extremely familiar with Fortran/C++
- Be prepared: don't plan to do this when you are highly pressed for time
- Doing this in Linux will typically save a lot of work
- I'd be happy to assist, but my abilities to do so are entirely experience based and include large gaps of knowledge
 - Like you, I'm not a computer scientist, and just want to get it running
- Just for this reason, it might be preferable to work with fmincon for this class

- Low-thrust trajectory optimization
- Direct transcription through Gauss-Lobatto collocation
- Some appealing aspects...
 - Very direct control over problem
 - Extremely sparse Jacobian
 - Embarrassingly parallel formulation
- Has proven to be very effective with large and/or complex problems



- Input state
- Control vector
- Propagated state
- . - Constraint vector
- ★ Target




→ This is all you directly control


● Propagated state

- - - Constraint vector


★ Target



● Input state → Control vector ● Propagated state - · - Constraint vector ★ Target




● Input state → Control vector ● Propagated state - · - Constraint vector ★ Target




• Input state → Control vector • Propagated state - - - Constraint vector ★ Target

- Outward spiral from GEO
 - ~5 days flight time
 - Two-body motion
 - 3000 kg initial mass
 - 50 kW constant power (60% jet efficiency)
 - 2000 s specific impulse
- Objective: maximize orbital energy

Initial guess



Solution






(Video)

Project suggestion

- It is actually quite manageable to write a simple but reasonably powerful low-thrust tool
 - The final product can be about 100-150 lines of code, plus some 50-100 lines to initialize the problem
 - Took me about an hour to write it from scratch (where I count several years of experience with these tools as “scratch”)
- What do you need?
 - A ready-to-go optimizer that can compute derivatives through finite differencing (fmincon?)
 - A propagator for two-body motion (MATLAB's rk4?)
 - Perhaps start by building the POPSICLE-problem to get the interface down, after which all you need to do is change the contents of the function

- Suggested problem
 - 180 degree Earth-Mars transfer
 - 10 trajectory segments, single shooter
 - 2000 kg initial mass
 - 30 kW of constant power
 - Fixed specific impulse (but can be optimized with the trajectory)
 - Flight time of 250 days
- Variables
 - Hardcode the initial & final state, including the initial mass
 - That leaves 1 variable for power, 1 variable for the specific impulse, and three thrust components for 10 segments: 32 variables
- Constraints
 - 3 for the final position, 3 for the final velocity, and 10 thrust constraints for a total of 16 constraints

- Propagation
 - Divide total flight time (250 days) by 10
 - Apply the 3 thrust components of the current segment as an impulsive Delta-V to the initial state (just set all of these to 10% of max thrust for the initial guess)
 - Store the thrust magnitude for the correct thrust constraint
 - Propagate for the 25 days of the current segment
 - Repeat, with as initial state the final state of the last 25 day propagation
 - At the end of the 10 segments, compute the error with respect to the target state (Mars) and store these as the appropriate constraint values
- Cost function
 - Take the mass at the end of the propagation, either maximize its positive value or minimize its negative value (most tools minimize by default)



- Perhaps a bit much to take in right now, but a very manageable exercise
- What you'll get out of it...
 - Familiarity with these types of tools
 - Insight into the effects of scaling
 - Optimizer will independently introduce 2 thrust arcs and 1 coast arc
 - Optimizer will independently optimize the value of the specific impulse and/or spacecraft power if you open up those variables
 - Once you've done this once in such a controlled formulation, you'll be very comfortable using these tools for much more advanced problems (you could actually evolve this simple example into a very powerful low-thrust optimization tool...)

- High-level overview of the key aspects & features of optimization packages
- Hopefully, this is a useful starting point for building your own IPOPT/SNOPT/fmincon/... implementations
 - Try out the POPSICLE-problem!
 - Slides & supporting material will be online
- Contact me if you need a hand with these packages and/or the problem suggestion!
 - jon.herman@colorado.edu

